

Univerzitet u Nišu  
Elektronski fakultet Niš

Seminarski rad

## **High availability primeru Redis baze podataka**

Sistemi za upravljanje bazama podataka 2019/20

Mentor:  
Aleksandar Stanimirović

Student:  
Todor Majstorović 1088

# Uvod

Često se može desiti pad performansi servera, poput naglog zagušenja ili nestanka struje. Da bi se takvi slučajevi sprečili, potrebno je osposobiti sistem da se izbori sa takvim problemima. Rešenje je korišćenje High Availability (HA) arhitekture, gde se servisi implementiraju tako da osiguraju optimalan rad servera. Iako ne postoje jedinstvena rešenja, postoje preporuke koje se trebaju ispoštovati.

## 1. Koristiti više servera

Ukoliko se koristi samo jedan server, pad tog servera dovodi do pada cele aplikacije dok se taj server ne restartuje.

Očigledno rešenje je deployovanje vaše aplikacije na više servera, vodeći računa o pravilnoj raspodeli zahteva na svaki od tih servera. Moguće je čak i rasporediti delove servera, poput dela koji procesira fajlove ili se bavi obradom mailova.

## 2. Skaliranje baza podataka

Baze podataka su najpopularniji način za čuvanje podataka. Podjednako su važne kao i sam server. Baze se pokreću na odvojenim serverima i moguće je da taj server otkáže, što može dovesti do gubljenja korisničkih podataka.

Redundansa je proces kreiranja sistema sa više nivoa dostupnosti koji može da detektuje otkaz i da se automatski oporavi. Ovo se postiže održavanjem slave baza, koje se ažuriraju zajedno sa glavnom bazom. Još jedan bitan koncept je shardovanje. Shard je horizontalna particija baze, gde se redovi tabela dele na više servera.

# Redis

Redis je open-source in-memory baza podataka koja je dizajnirana da bude brza i jednostavna. Omogućava real-time komunikaciju, većina upita se izvršavaju kraće od milisekunde što dopušta veliki broj konkurentnih upita. Kako bi osigurao visoku dostupnost, oporavak od grešaka i replikaciju, koriste se Redis cluster i Redis Sentinel.

Instanca Redisa predstavlja proces koji se ponaša kao in-memory skladište. Svi upiti za čitanje i upis idu ka tom procesu. Kako bi se zaštitili od otkaza, Redis može da replicira podatke, kao i da šalje promene većem broju instanca ili replika.

Redis koristi unidirekcionu asinhronu replikaciju od mastera ka slave-ovima. Jedan master može imati veći broj replika, a replike mogu imati svoje replike. Kako bi se osiguralo što kraće kašnjenje, replikacije je neblokirajuća i izvršava se nakon što se izvrši na masteru. Moguće je da se izgube podaci ako se otkaz desi između izvršenja operacije i izvršenja replikacije. Jedno rešenje je da se čeka dok se ne napravi odredjen broj replika.

Redis replike mogu biti dodavane i konfigurisane u runtime, a podržana je i parcijalna replikacija. Replike mogu da se podese da služe kao hot-standby, da opslužuju samo upite za čitanje, a mogu čak i da dozvole upise.

## Pros:

- Lako i jednostavno za podešavanje
- Može biti automatizovano
- Nastavlja da radi dok je jedna instanca mastera dostupna

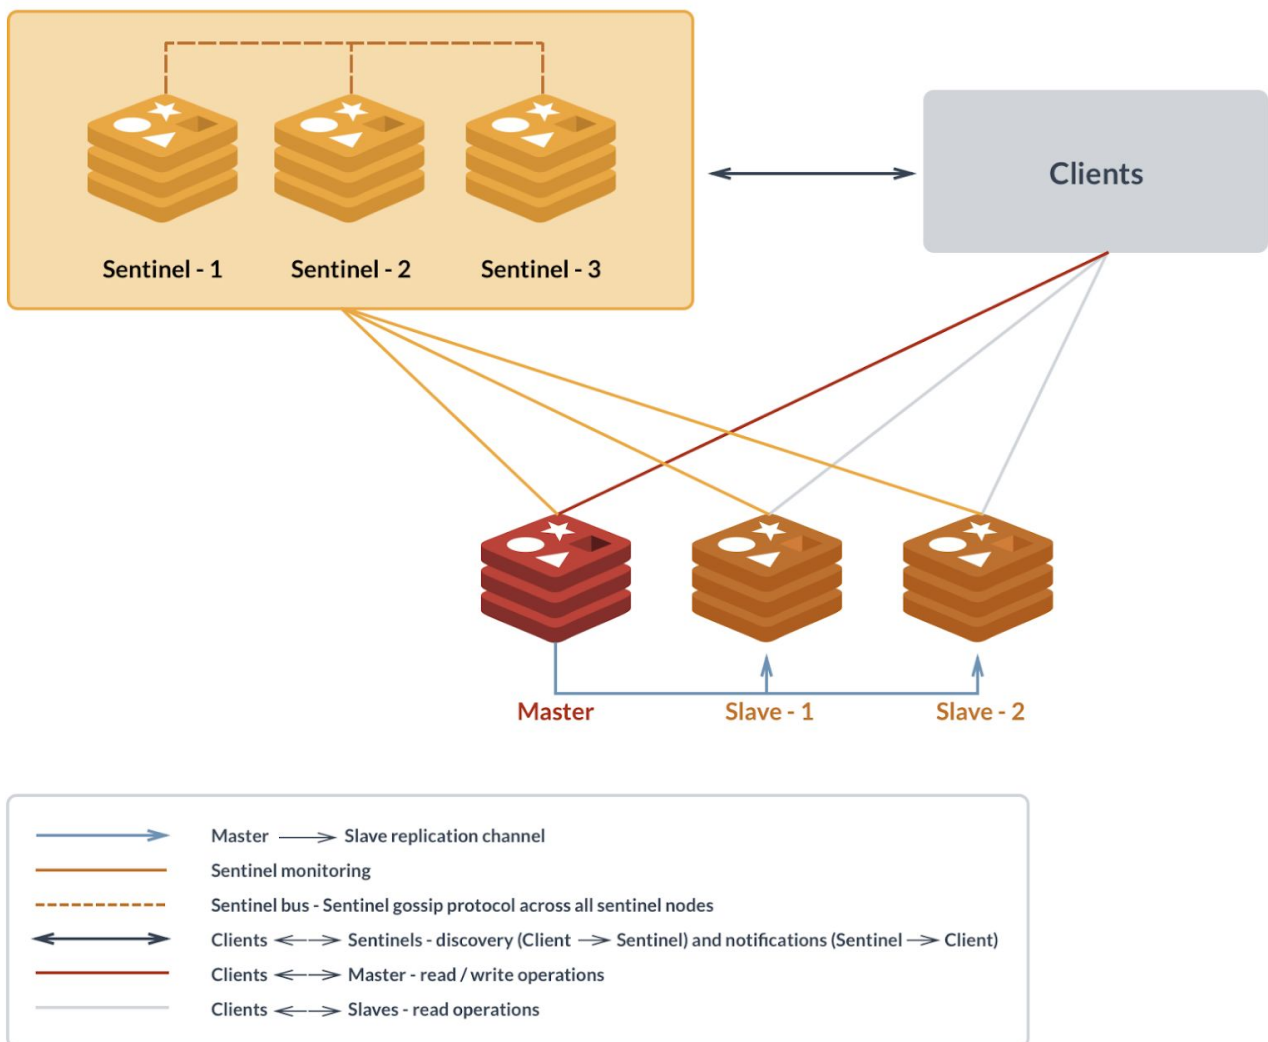
## Cons:

- Upisi moraju da idu na master
- Slave može da opslužuje čitanje, ali mogu se dobiti nevalidni podaci
- Ne sharduje podatke
- Ako otkáže master, ručno se mora unaprediti novi

Problem kod Redis replika je što u slučaju otkaza mastera mora ručno da se unapredi slave u novi master. Ovaj problem se rešava korišćenjem Redis Sentinela.

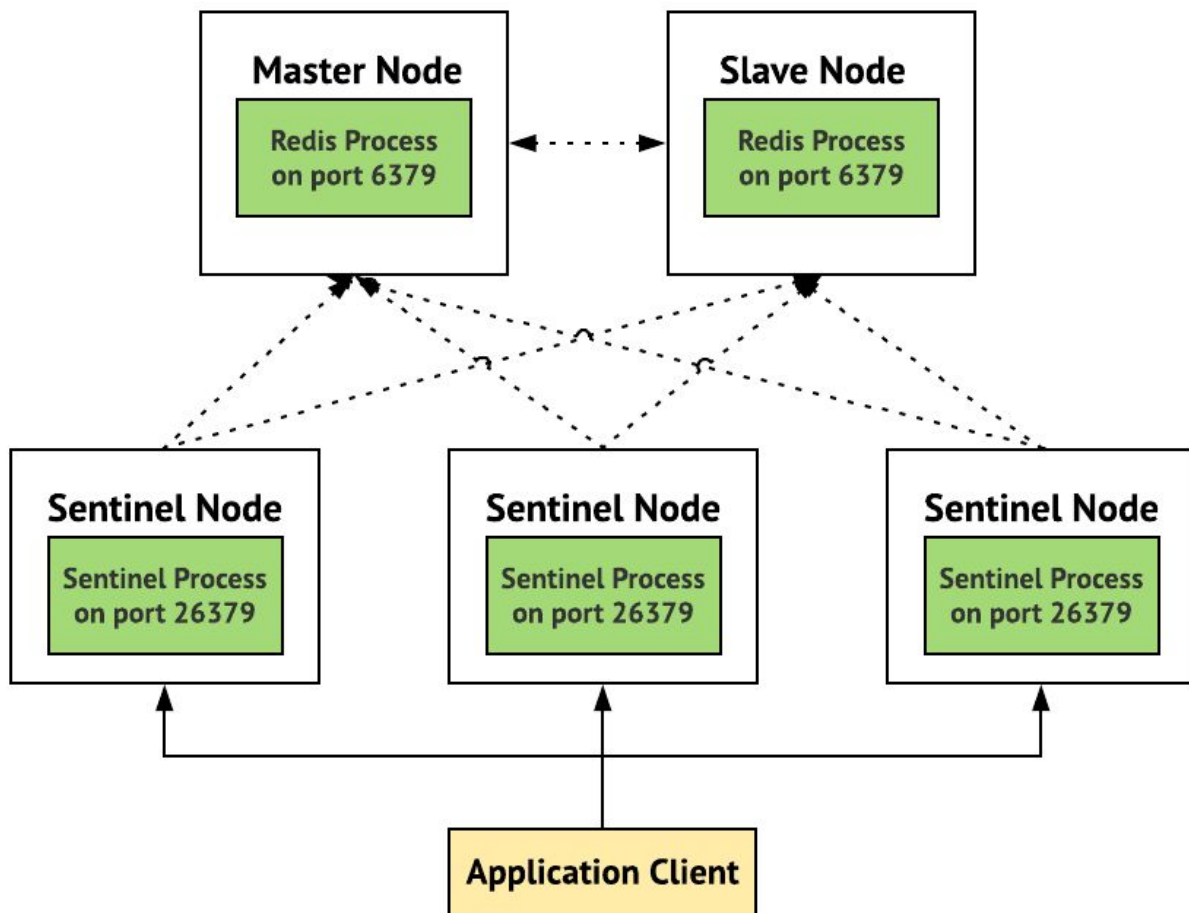
# Redis Sentinel

Kada se koristi Redis sa jednom instancom, gde jedan Redis server opslužuje celu bazu, Redis Sentinel se bavi dostupnošću. Sentinel je komponenta Redis servera koja se mora deployovati pored servera i replika. Sentinel se zasniva na kvorumu kako bi otkrio otkaz, pa zahteva više instanci koje će se nadgledati. Kada se većina Sentinelova složi da je otkaz nastupio, jedna od replika postaje master. Redis Sentinel pruža i interfejs za otkrivanje servisa koje klijent koristi da se poveže sa glavnom Redis instancom.



Kada Sentinel detektuje otkaz, jedna od replika postaje master. Svi ostali slave-ovi se automatski konfiguriraju da koriste novi master. Kada klijent postavi upit na Sentinel, Sentinel mu šalje adresu mastera.

Kako Sentinel izvršava automatski oporavak:



**Configuration: quorum = 2**

Quorum vrednost predstavlja broj sentinela koji treba da potvrde otkaz. Medjutim to nije dovoljno, jedan od sentinela mora da se proglasi vodjom i da se autorizuje.

# Podešavanje replika i Sentinela

Pokrenuti master na portu 6379

```
4385:M 22 Jun 2020 22:31:40.969 # Server initialized
4385:M 22 Jun 2020 22:31:40.969 * Ready to accept connections

4385:M 22 Jun 2020 22:32:04.795 * Replica 127.0.0.1:6380 asks for synchronization
4385:M 22 Jun 2020 22:32:04.795 * Full resync requested by replica 127.0.0.1:6380
4385:M 22 Jun 2020 22:32:04.795 * Replication backlog created, my new replication
IDs are 'd6b00110ecb62b878b462d1d0b3ec8d3192f7224' and '00000000000000000000000000000000'
4385:M 22 Jun 2020 22:32:04.795 * Starting BGSAVE for SYNC with target: disk
4385:M 22 Jun 2020 22:32:04.796 * Background saving started by pid 4396
4396:C 22 Jun 2020 22:32:04.797 * DB saved on disk
4385:M 22 Jun 2020 22:32:04.855 * Background saving terminated with success
4385:M 22 Jun 2020 22:32:04.855 * Synchronization with replica 127.0.0.1:6380 succeeded
```

Pokrenuti repliku komandom :

redis-server --port 6380 --slaveof 127.0.0.1 6379

```
4395:S 22 Jun 2020 22:32:04.794 # Server initialized
4395:S 22 Jun 2020 22:32:04.794 * Ready to accept connections
4395:S 22 Jun 2020 22:32:04.794 * Connecting to MASTER 127.0.0.1:6379
4395:S 22 Jun 2020 22:32:04.795 * MASTER <-> REPLICHA sync started
4395:S 22 Jun 2020 22:32:04.795 * Non blocking connect for SYNC fired the event.
4395:S 22 Jun 2020 22:32:04.795 * Master replied to PING, replication can continue
...
4395:S 22 Jun 2020 22:32:04.795 * Partial resynchronization not possible (no cache
d master)
4395:S 22 Jun 2020 22:32:04.796 * Full resync from master: d6b00110ecb62b878b462d1
d0b3ec8d3192f7224:0
4395:S 22 Jun 2020 22:32:04.855 * MASTER <-> REPLICHA sync: receiving 175 bytes from master to disk
4395:S 22 Jun 2020 22:32:04.855 * MASTER <-> REPLICHA sync: Flushing old data
4395:S 22 Jun 2020 22:32:04.855 * MASTER <-> REPLICHA sync: Loading DB in memory
4395:S 22 Jun 2020 22:32:04.855 * Loading RDB produced by version 6.0.5
4395:S 22 Jun 2020 22:32:04.855 * RDB age 0 seconds
4395:S 22 Jun 2020 22:32:04.855 * RDB memory usage when created 2.01 Mb
4395:S 22 Jun 2020 22:32:04.855 * MASTER <-> REPLICHA sync: Finished with success
```

```
#Master 1 sentinel 1
```

```
sentinel monitor mymaster 127.0.0.1 6379 2
```

```
#Replika 1 sentinel 2
```

```
sentinel monitor mymaster 127.0.0.1 6380 2
```

```
#Replika 1 sentinel 3
```

```
sentinel monitor mymaster 127.0.0.1 6380 2
```

Podesite port sentinela i adresu koju će sentinel da nadgleda

```
#  
# Note: master name should not include special characters or spaces.  
# The valid charset is A-z 0-9 and the three characters ".-_  
sentinel monitor mymaster 127.0.0.1 6379 2  
  
# sentinel auth-pass <master-name> <password>  
#  
# Set the password to use to authenticate with the master and replicas.  
# Useful if there is a password set in the Redis instances to monitor.  
#  
# Note that the master password is also used for replicas, so it is not  
# possible to set a different password in masters and replicas instances  
# if you want to be able to monitor these instances with Sentinel.  
#  
# However you can have Redis instances without the authentication enabled  
# mixed with Redis instances requiring the authentication (as long as the  
# password set is the same for all the instances requiring the password) as  
# the AUTH command will have no effect in Redis instances with authentication  
# switched off.  
#  
# Example:
```

Broj 2 predstavlja podešavanje kvoruma

```

GNU nano 2.0.6      File: /usr/local/etc/redis-sentinel.conf      Modified
# In the Redis servers side, the ACL to provide just minimal access to
# Sentinel instances, should be configured along the following lines:
#
#   user sentinel-user >somepassword +client +subscribe +publish \
#                               +ping +info +multi +slaveof +config +client +exec on
#
# sentinel down-after-milliseconds <master-name> <milliseconds>
#
# Number of milliseconds the master (or any attached replica or sentinel) should
# be unreachable (as in, not acceptable reply to PING, continuously, for the
# specified period) in order to consider it in S_DOWN state (Subjectively
# Down).
#
# Default is 30 seconds.
sentinel down-after-milliseconds mymaster 5000

# requirepass <password>
#
# You can configure Sentinel itself to require a password, however when doing
# so Sentinel will try to authenticate with the same password to all the
# other Sentinels. So you need to configure all your Sentinels in a given

```

Podesite vreme koliko treba da prodje da se promeni stanje mastera nakon što padne.

```

#   to a Sentinel current configuration, to be forced to replicate
#   with the right master, is exactly the failover timeout (counting since
#   the moment a Sentinel detected the misconfiguration).
#
# - The time needed to cancel a failover that is already in progress but
#   did not produced any configuration change (SLAVEOF NO ONE yet not
#   acknowledged by the promoted replica).
#
# - The maximum time a failover in progress waits for all the replicas to be
#   reconfigured as replicas of the new master. However even after this time
#   the replicas will be reconfigured by the Sentinels anyway, but not with
#   the exact parallel-syncs progression as specified.
#
# Default is 3 minutes.
sentinel failover-timeout mymaster 180000

# SCRIPTS EXECUTION
#
# sentinel notification-script and sentinel reconfig-script are used in order
# to configure scripts that are called to notify the system administrator

```

<sup>^</sup>G Get Help    <sup>^</sup>O WriteOut    <sup>^</sup>R Read File    <sup>^</sup>Y Prev Page    <sup>^</sup>K Cut Text    <sup>^</sup>C Cur Pos  
<sup>^</sup>X Exit        <sup>^</sup>J Justify    <sup>^</sup>W Where Is    <sup>^</sup>V Next Page    <sup>^</sup>U UnCut Text <sup>^</sup>T To Spell



Podesite broj replika koje se trebaju rekonfigurisati da koriste novi master nakon failovera, u našem slučaju pošto imamo 2 replike, 1 je potrebno podesiti

```
# group with the same "requirepass" password. Check the following documentation
# for more info: https://redis.io/topics/sentinel

# sentinel parallel-syncs <master-name> <numreplicas>
#
# How many replicas we can reconfigure to point to the new replica simultaneously
# during the failover. Use a low number if you use the replicas to serve query
# to avoid that all the replicas will be unreachable at about the same
# time while performing the synchronization with the master.
sentinel parallel-syncs mymaster 1

# sentinel failover-timeout <master-name> <milliseconds>
#
# Specifies the failover timeout in milliseconds. It is used in many ways:
#
# - The time needed to re-start a failover after a previous failover was
#   already tried against the same master by a given Sentinel, is two
#   times the failover timeout.
#
# - The time needed for a replica replicating to a wrong master according

^G Get Help  ^O WriteOut  ^R Read File ^Y Prev Page ^K Cut Text  ^C Cur Pos
^X Exit      ^J Justify   ^W Where Is ^V Next Page ^U UnCut Text ^T To Spell
```

Restartujte Sentinel servis na svim čvorovima.

```
systemctl restart redis-sentinel
```

Restartujte Sentinel servis na svim čvorovima i otvorite portove kako bi Sentinel mogao da prima poruke na svom portu

```
firewall-cmd --zone=public --permanent
--add-port=26379/tcp
firewall-cmd --reload
```

Sve replike ce biti automatski otkrivene nakon pokretanja sve tri instance sentinela.

```
(it was originally set to 256).

Redis 6.0.5 (1f24f59f/0) 64 bit
Running in sentinel mode
Port: 5000
PID: 4480

http://redis.io

4480:X 22 Jun 2020 22:38:59.284 # Sentinel ID is d9c8e23bb9e3516bde021b0c71862022b0a98458
4480:X 22 Jun 2020 22:38:59.284 # +monitor master mymaster 127.0.0.1 6379 quorum 2
```

## Sentinel 1

```
4485:X 22 Jun 2020 22:39:21.010 * Increased maximum number of open files to 10032
(it was originally set to 256).

Redis 6.0.5 (1f24f59f/0) 64 bit
Running in sentinel mode
Port: 5001
PID: 4485

http://redis.io

4485:X 22 Jun 2020 22:39:21.011 # Sentinel ID is d9c8e23bb9e3516bde021b0c71862022b0a98458
4485:X 22 Jun 2020 22:39:21.011 # +monitor master mymaster 127.0.0.1 6380 quorum 2
```

## Sentinel 2

## Sentinel 3

## Testiranje Failovera

```
127.0.0.1:5000> SENTINEL get-master-addr-by-name mymaster
1) "127.0.0.1"
2) "6380"
```

# Zaključak

U radu je opisan jedan od problema koji se tiče svih baza podataka, a i servera uopšte, a to je problem visoke dostupnosti prilikom otkaza. Najbolje rešenje je korišćenje više servera kako bi se moglo preuzeti u slučaju otkaza jednog od njih. Predstavljeno je rešenje na primeru Redis baze podataka, gde se za automatsko pokretanje procesa oporavka koristi Redis sentinel. Dat je primer konfigurisanja redis instanci i redis sentinela koje nadgledaju napravljene mastere i replike. Na kraju je dat i primer oporavka

## Reference

1. <https://medium.com/@amila922/redis-sentinel-high-availability-everything-you-need-to-know-from-dev-to-prod-complete-guide-deb198e70ea6>
2. <https://www.tecmint.com/setup-redis-high-availability-with-sentinel-in-centos-8/>
3. <https://redis.io/topics/sentinel>
4. <https://medium.com/@petehouston/install-and-config-redis-on-mac-os-x-via-homebrew-eb8df9a4f298>